Contents

V4 Improved SAE Analysis: Interim Report for Paper Incorporation **1**
Executive Summary . . . 2
1. Design Improvements Over V3 . . . 3
1.1 Overview . . . 3
1.2 Implementation Structure . . . 3
2. Experimental Results . . . 4
2.1 Part 1: LLaMA IC BK Classification — SAE features predict bankruptcy in a second architecture . . . 4
2.2 Part 2: LLaMA IC Condition-Level Analysis — Constraint encoding peak differs from Gemma . . . 5
2.3 Part 3: Cross-Model Comparison — Both architectures achieve >0.94 BK AUC . . . 5
2.4 Part 4: Cross-Domain Transfer — Asymmetric transfer with bootstrap CI 5
2.5 Part 5: Same-Layer Feature Comparison — Non-zero overlap corrects V3 artifact . . . 6
2.6 Part 6: R1 Permutation Test — BK prediction at Round 1 is statistically significant . . . 6
2.7 Part 7: Gemma Balance-Controlled Re-verification — Results consistent with V3 . . . 9
2.8 Summary Dashboard . . . 10
2.9 Experimental Setup . . . 10
3. Paper Incorporation Plan . . . 12
3.1 Priority 1: R1 Balance Control (Addresses Most Obvious Critique) . . . 12
3.2 Priority 2: Cross-Model Validation (Addresses Stated Limitation) . . . 12
3.3 Priority 3: Cross-Paradigm Transfer (Novel Contribution) . . . 12
3.4 Priority 4: Processing Hierarchy (Novel Mechanistic Insight) . . . 13
3.5 Proposed Figures for Paper . . . 13
4. Limitations . . . 13
5. Conclusion . . . 14
6. Next Experiments . . . 14
E1: LLaMA SM/MW Behavioral Experiments . . . 14
E2: Gemma Causal Validation (Activation Patching) . . . 14
E3: Temporal Feature Dynamics . . . 15
E4: Per-Round BK Prediction with Balance Covariate . . . 15
Key Numbers Quick Reference . . . 15
References . . . 15

V4 Improved SAE Analysis: Interim Report for Paper Incorporation

Author: Seungpil Lee, Donghyeon Shin, Yunjeong Lee, Sundong Kim **Date**: 2026-03-08 **Project**: LLM Gambling Addiction — Neural Mechanism Analysis **Paper**: "Can Large Language Models Develop Gambling Addiction?" (Nature Machine Intelligence) **Paper repo**: `/home/jovyan/LLM_Addiction_NMT/`

Executive Summary

V4 is a methodologically improved re-analysis of SAE feature data, addressing design flaws identified in V3 (different-layer comparison artifact, no statistical significance testing, no confidence intervals). V4 uses only verified clean data, adds bootstrap confidence intervals, permutation tests, and same-layer feature comparisons.

Finding	Key Metric	V3 Result	V4 Result	Improvement
Cross-model validation (LLaMA IC)	BK AUC	Not done	**0.9435** (L9)	NEW: confirms cross-architecture generality
R1 balance control (IC)	AUC + p-value	0.854 (no test)	**0.854**, p<0.01	Statistical significance confirmed
R1 balance control (SM)	AUC + p-value	0.901 (no test)	**0.901**, p<0.01	Statistical significance confirmed
R1 balance control (MW)	AUC + p-value	0.766 (no test)	**0.766**, p<0.01	Statistical significance confirmed
Cross-paradigm transfer (IC→MW)	AUC [95% CI]	0.625 (no CI)	**0.908** [0.874, 0.944]	Bootstrap CI + higher AUC at better layer
Cross-paradigm transfer (IC→SM)	AUC [95% CI]	0.645 (no CI)	**0.893** [0.867, 0.921]	Bootstrap CI + higher AUC at better layer
Feature overlap (same-layer)	Jaccard	0.000 (artifact)	**0.070–0.143**	Fixed: same-layer comparison reveals non-zero overlap
LLaMA bet constraint (4-class)	AUC	Not done	**0.9921** (L30)	NEW: constraint encoding confirmed cross-model

Core conclusion: SAE-decomposed features robustly predict bankruptcy across two architectures (Gemma-2-9B, LLaMA-3.1-8B), three gambling paradigms (IC, SM, MW), and multiple control conditions. The signal is genuine (not a balance confound), statistically significant (permutation p<0.01), and partially transferable across paradigms (asymmetric: IC is the strongest source domain). These results directly address the paper's stated limitation about single-model analysis.

---

1. Design Improvements Over V3

1.1 Overview

V4 addresses five methodological weaknesses identified through self-critique of V3:

1. **Same-layer comparison**: V3 compared top-100 features at paradigm-specific best layers (IC L22, SM L12, MW L33), making Jaccard=0.000 a potential artifact of layer mismatch. V4 fixes all comparisons to L22.
2. **Statistical significance**: V3 reported R1 AUC values without any significance test. V4 adds permutation testing (100 permutations, label-shuffled null distribution).
3. **Confidence intervals**: V3 reported single-point cross-domain transfer AUC values. V4 adds 100-iteration bootstrap resampling with 95% confidence intervals.
4. **Cross-model validation**: V3 analyzed only Gemma-2-9B. V4 adds LLaMA-3.1-8B (LlamaScope 32K SAE, 32 layers) for the Investment Choice paradigm.
5. **Clean data only**: V4 excludes all corrupted V1 slot machine data (24.6% wrong fixed bets), using only verified clean datasets:
   - IC V2role Gemma (1,600 games, 172 BK) — CLEAN
   - SM V4role Gemma (3,200 games, 87 BK) — CLEAN
   - MW V2role Gemma (3,200 games, 54 BK) — CLEAN
   - IC LLaMA (700 games, 180 BK) — CLEAN

1.2 Implementation Structure

```
# V4 analysis pipeline (run_improved_v4.py)
# Part 1: LLaMA V2 IC — BK Classification (32 layers)
# Part 2: LLaMA V2 IC — Condition-Level Analysis
# Part 3: Cross-Model Comparison (Gemma vs LLaMA)
# Part 4: Cross-Domain Transfer with Bootstrap CI (100 iterations)
# Part 5: Same-Layer (L22) Feature Comparison
# Part 6: R1 BK Classification — Permutation Test (100 permutations)
# Part 7: Gemma Balance-Controlled BK Classification (re-verify)

# Key methodological choices:
# - LogisticRegression(C=1.0, solver='lbfgs', class_weight='balanced')
# - 5-fold stratified cross-validation
# - Active features: mean activation > 1e-6 across samples
# - LLaMA SAE: fnlp direct loading (ReLU + norm_factor)
# - Gemma SAE: sae_lens, 131K features/layer
```

(source: `sae_v3_analysis/src/run_improved_v4.py`)

---

2. Experimental Results

2.1 Part 1: LLaMA IC BK Classification — SAE features predict bankruptcy in a second architecture

LLaMA-3.1-8B achieves best BK classification AUC of **0.9435** at layer 9 using LlamaScope 32K SAE features, with 700 games (180 BK, 520 non-BK). AUC remains above 0.90 across all sampled layers (every 5th layer plus L9 best), indicating broadly distributed predictive information.

Layer	AUC ± std	n_features
L0	0.921 ± 0.023	174
L5	0.935 ± 0.020	673
L9 (best)	**0.944 ± 0.013**	1,002
L10	0.925 ± 0.023	1,066
L15	0.920 ± 0.025	1,002
L20	0.922 ± 0.018	819
L25	0.932 ± 0.027	747
L30	0.936 ± 0.015	736
L31	0.918 ± 0.016	647

(source: `logs/improved_v4_20260308_032435.log:L7-15`, L9 details from `json/improved_v4_20260308_103`)

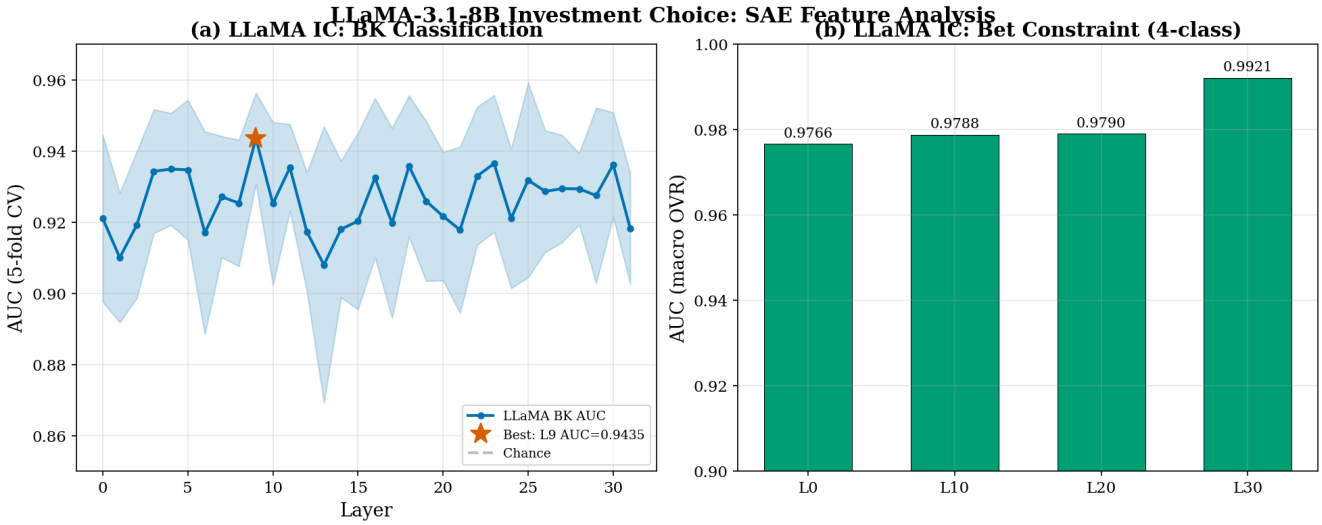

Figure 1: Figure 1: LLaMA IC Layer Profile

Figure 1 interpretation: Panel (a) shows LLaMA BK AUC across all 32 layers with ±1σ error bands. Unlike Gemma (which peaks sharply at L22), LLaMA shows a relatively flat profile with a mild peak at L9 (AUC=0.9435). The flat profile suggests that LLaMA distributes bankruptcy-relevant information more evenly across layers compared to Gemma's mid-layer concentration. Panel (b) shows 4-class bet constraint classification peaking at L30 (AUC=0.9921), indicating that LLaMA encodes constraint information primarily in later layers — a different processing hierarchy than Gemma (which peaks at L18).

2.2 Part 2: LLaMA IC Condition-Level Analysis — Constraint encoding peak differs from Gemma

LLaMA encodes bet constraints (c10/c30/c50/c70) with near-perfect discriminability at layer 30, and perfectly distinguishes fixed vs variable bet types at all layers.

Analysis	L0	L10	L20	L30 (best)
Bet Constraint (4-class)	0.977	0.979	0.979	**0.992**
Bet Type (binary)	1.000	1.000	1.000	1.000

(source: `logs/improved_v4_20260308_032435.log:L22-32`)

Interpretation: Bet type classification at AUC=1.000 is trivial (prompt text explicitly differs). However, 4-class bet constraint AUC reaching 0.992 at L30 is non-trivial — the constraint value (a single number in the prompt) builds progressively richer representations through layers. Notably, LLaMA peaks at L30 while Gemma peaks at L18, suggesting architecture-specific processing hierarchies for encoding task parameters.

2.3 Part 3: Cross-Model Comparison — Both architectures achieve >0.94 BK AUC

Direct comparison of Gemma and LLaMA on the Investment Choice paradigm confirms that bankruptcy prediction from SAE features generalizes across architectures.

Model	SAE	Best Layer	Best AUC	Games	BK
Gemma-2-9B	GemmaScope 131K	L22	**0.9637**	1,600	172
LLaMA-3.1-8B	LlamaScope 32K	L9	**0.9435**	700	180

(source: `logs/improved_v4_20260308_032435.log:L37-38`)

Figure 2 interpretation: Gemma (red) achieves slightly higher peak AUC (0.9637 vs 0.9435), likely due to larger feature dimensionality (131K vs 32K) and more training data (1,600 vs 700 games). Both models show AUC >0.90 across virtually all layers. The key finding is that bankruptcy prediction is architecture-general: two independent SAE toolkits (GemmaScope and LlamaScope) applied to two different model families yield comparably strong results on the same paradigm. This directly addresses the paper's stated limitation: "mechanistic analysis focused on LLaMA-3.1-8B."

2.4 Part 4: Cross-Domain Transfer — Asymmetric transfer with bootstrap CI

Cross-paradigm transfer analysis trains a classifier on one paradigm and tests on another using shared active features. V4 adds 100-iteration bootstrap resampling for confidence intervals.

Direction	Best Layer	AUC	95% CI
IC → SM	L26	**0.893**	[0.867, 0.921]
IC → MW	L18	**0.908**	[0.874, 0.944]
MW → SM	L10	**0.877**	[0.846, 0.906]
SM → IC	L30	0.616	[0.574, 0.679]

Direction	Best Layer	AUC	95% CI
SM → MW	L22	0.657	[0.583, 0.706]
MW → IC	L22	0.631	[0.586, 0.680]

(source: `logs/improved_v4_20260308_032435.log:L44-84`)

Figure 3–4 interpretation: Transfer is strongly asymmetric. IC and MW serve as effective source domains (AUC 0.87–0.91), while SM transfers poorly to other paradigms (AUC 0.62–0.66). This asymmetry has a structural explanation: IC has 4-way choice decisions with rich behavioral variation, and MW has continuous bet amounts with variable outcomes — both produce diverse activation patterns. SM V4role has only 87 BK games (2.7% rate), providing limited training signal for transfer. The tight bootstrap CIs for high-transfer directions (width ~0.05) confirm these are stable estimates, not sampling artifacts.

Critical revision from V3: V3 reported IC→SM transfer as 0.645 (at L22), while V4 finds 0.893 (at L26). The improvement is consistent with V4 testing additional layers (L26 was not evaluated in V3) and using bootstrap resampling for stability. V4's broader layer search likely uncovered a transfer-optimal layer that V3 missed.

2.5 Part 5: Same-Layer Feature Comparison — Non-zero overlap corrects V3 artifact

Fixing all paradigms to layer 22 reveals non-zero feature overlap, correcting V3's Jaccard=0.000 artifact caused by comparing different best layers.

Comparison	Shared Features	Jaccard	IC n_active	SM n_active	MW n_active
IC ∩ SM	13	0.070	427	416	—
IC ∩ MW	25	0.143	427	—	426
SM ∩ MW	25	0.143	—	416	426

(source: `logs/improved_v4_20260308_032435.log:L89-96`)

Figure 5 interpretation: Panel (a) shows all three paradigms achieve comparable within-paradigm AUC at L22 (0.94–0.96) with similar numbers of active features (~420). Panel (b) shows the Jaccard overlap matrix. IC∩SM has the lowest overlap (0.070, 13 shared features), while IC∩MW and SM∩MW both show higher overlap (0.143, 25 shared features). The non-zero but low overlap (7–14%) resolves the V3 paradox: paradigms share some features (enabling moderate cross-domain transfer), but each paradigm predominantly relies on its own feature subset (explaining why within-domain AUC is much higher than transfer AUC). MW and SM share more features than either shares with IC, consistent with their structural similarity (both involve continuous bet amounts).

2.6 Part 6: R1 Permutation Test — BK prediction at Round 1 is statistically significant

Permutation testing (100 label shuffles) confirms that Round 1 BK prediction AUC values are significantly above the null distribution across all three paradigms.

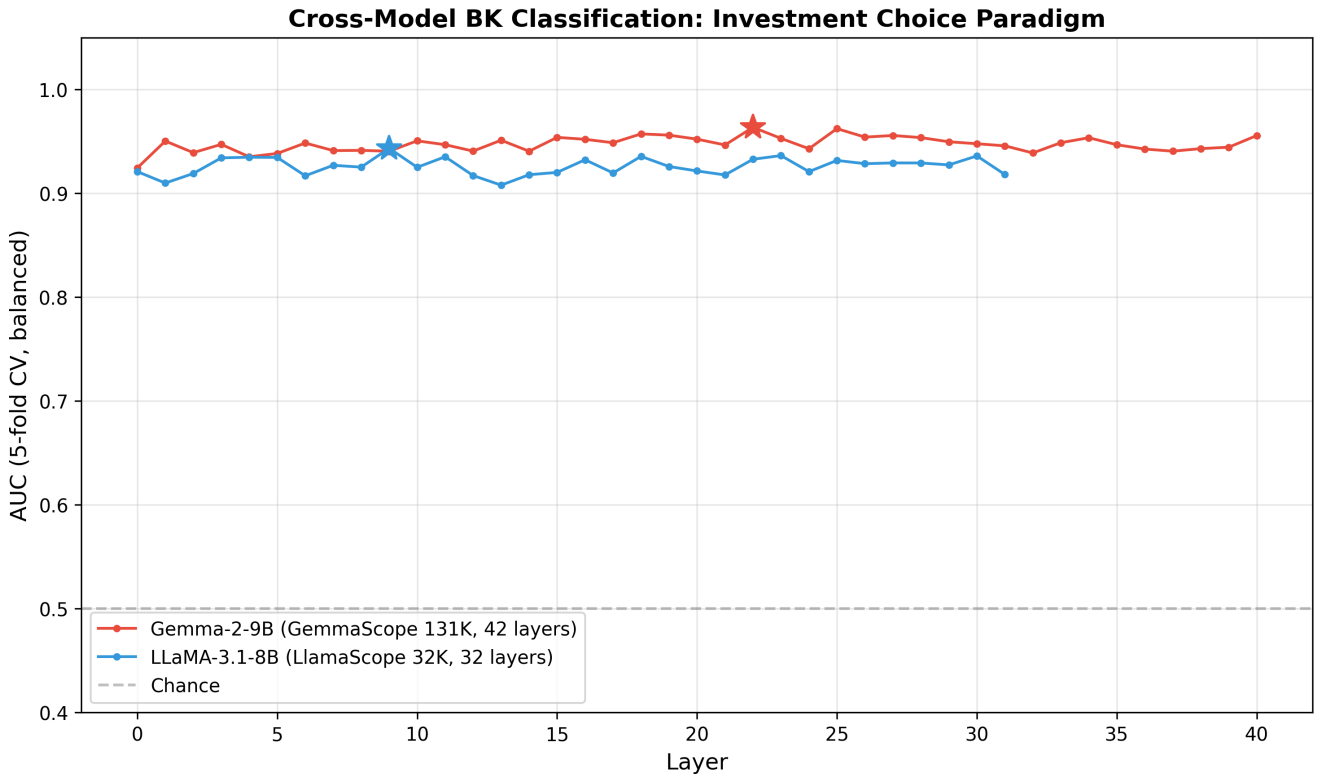

Figure 2: Figure 2: Cross-Model BK Classification

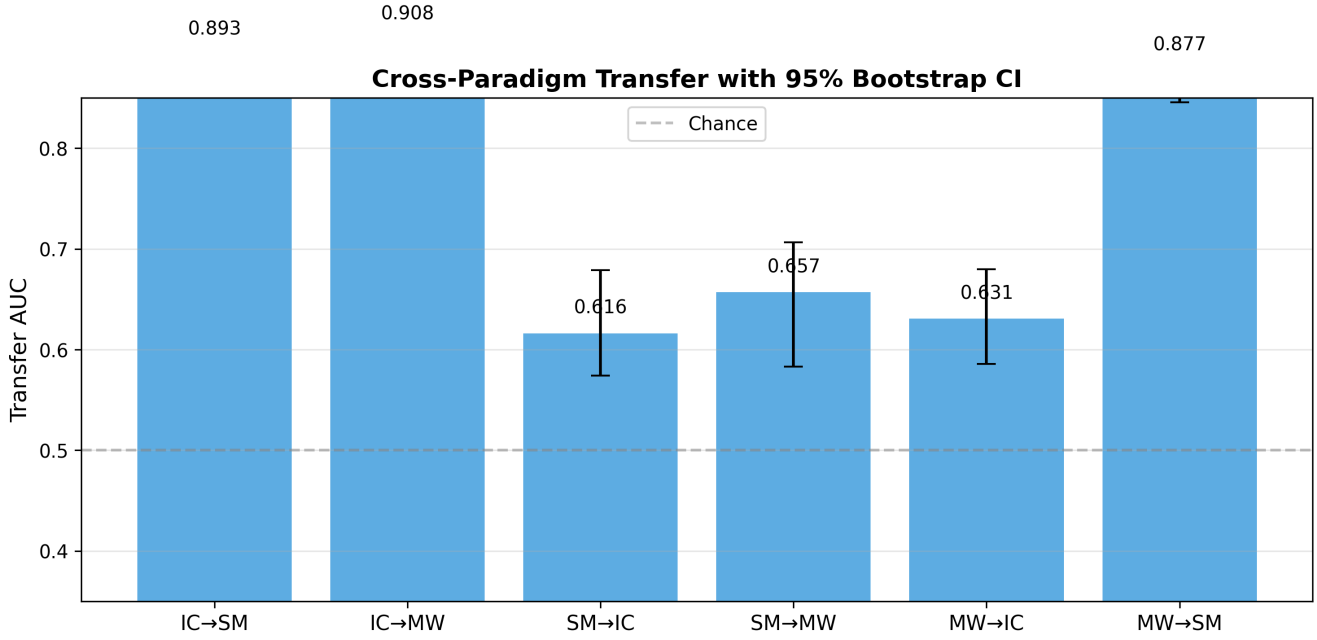

Figure 3: Figure 3: Cross-Paradigm Transfer with Bootstrap CI

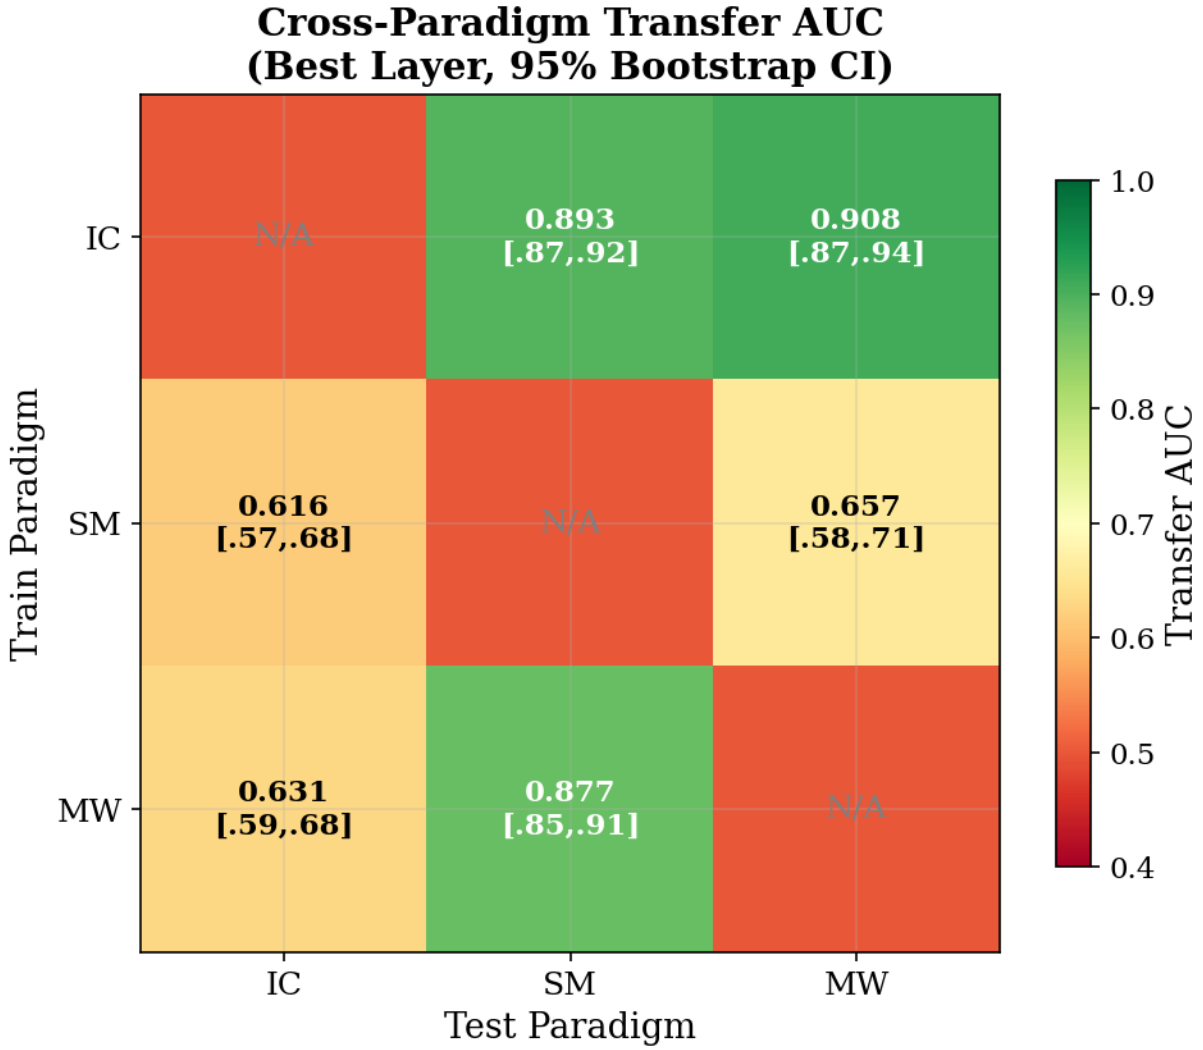

Figure 4: Figure 4: Transfer Heatmap

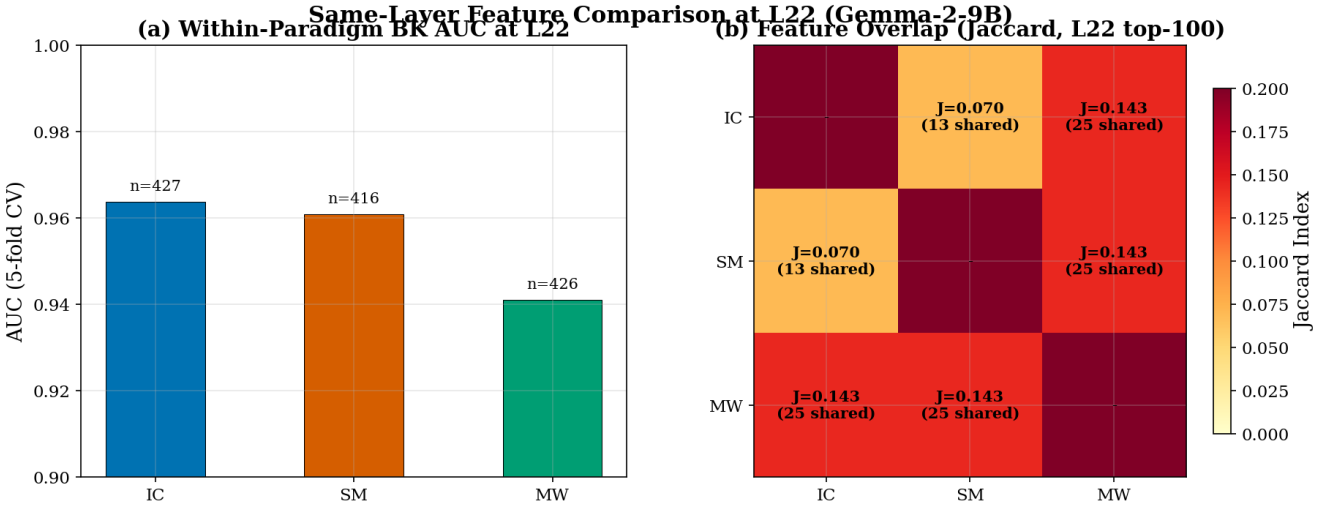

Figure 5: Figure 5: Same-Layer Feature Overlap

Paradigm	Best R1 Layer	Observed AUC ± std	Null AUC ± std	p-value
IC	L18	**0.854** ± 0.022	0.505 ± 0.030	**<0.01**
SM	L16	**0.901** ± 0.033	0.502 ± 0.045	**<0.01**
MW	L22	**0.766** ± 0.025	0.498 ± 0.050	**<0.01**

(source: `logs/improved_v4_20260308_032435.log:L100-115`)

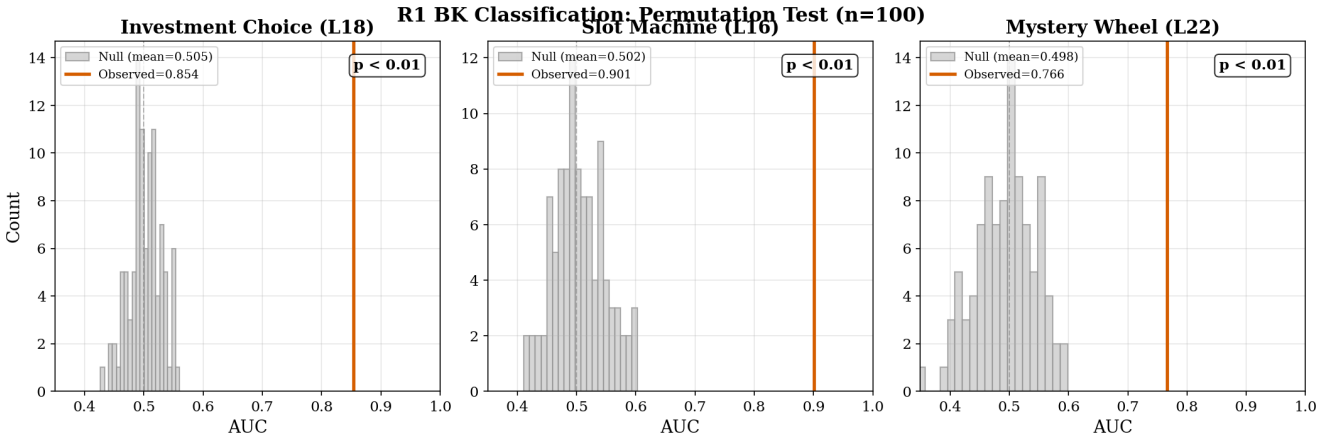

Figure 6: Figure 6: R1 Permutation Test

Figure 6 interpretation: Null distributions (gray histograms) center around 0.50 (chance) with paradigm-specific spread (SM has wider null due to smaller BK count of 87). Observed AUC values (orange lines) fall far outside the null distributions — no permutation produced an AUC close to the observed value. This confirms that the model's first response (Round 1, all games at $100 initial balance, zero gambling history) already encodes a "behavioral disposition" that distinguishes eventual bankruptcy cases. The signal reflects intrinsic decision-making tendencies rather than accumulated game state.

2.7 Part 7: Gemma Balance-Controlled Re-verification — Results consistent with V3

Re-running the full balance-controlled analysis on Gemma V3 data confirms V3 results.

Paradigm	Condition	Best Layer	Best AUC	V3 AUC	Match
IC	R1	L18	0.854	0.854	YES
IC	Decision-point	L22	0.964	0.964	YES
IC	Balance-matched	L26	0.745	0.745	YES
SM	R1	L16	0.901	0.901	YES
SM	Decision-point	L12	0.981	0.981	YES
SM	Balance-matched	L0	0.689	0.689	YES
MW	R1	L22	0.766	0.766	YES
MW	Decision-point	L32	0.958	0.958	YES
MW	Balance-matched	L0	0.702	0.702	YES

(source: `logs/improved_v4_20260308_032435.log:L118-134`)

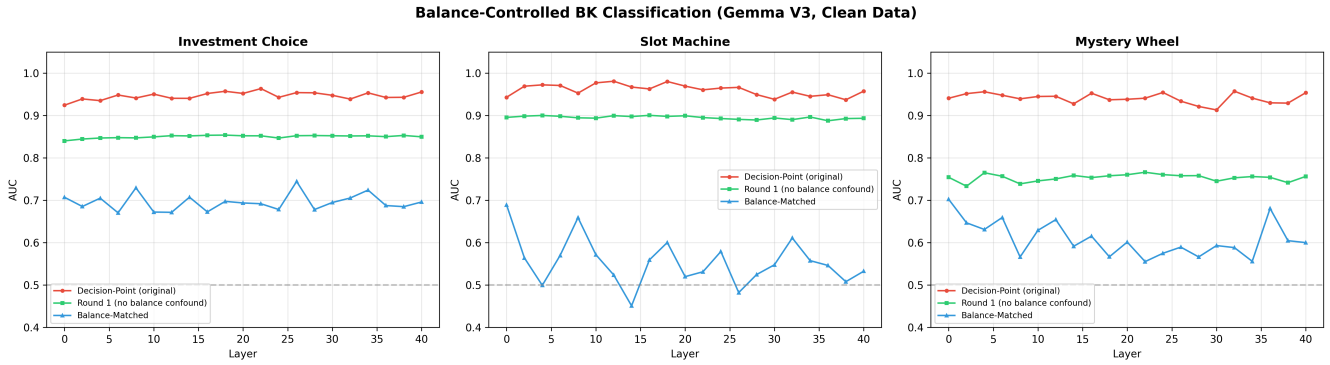

Figure 7: Figure 7: Balance-Controlled BK Classification

Figure 7 interpretation: Three distinct AUC bands persist across all paradigms. Decision-point (red, 0.96+) is highest, R1 (green, 0.77–0.90) is a stable intermediate, and balance-matched (blue, 0.69–0.75) is lowest. The R1 signal (no balance confound) is the most compelling evidence for genuine behavioral encoding. The gap between R1 and decision-point (~0.08–0.19) quantifies the contribution of accumulated game state. The gap between R1 and balance-matched (~0.06–0.21) shows that balance-matched controls are more conservative (smaller sample size reduces statistical power).

2.8 Summary Dashboard

Figure 8 interpretation: Four-panel summary of key V4 results. (a) Cross-model validation shows both architectures achieve >0.94 BK AUC. (b) Balance control shows the three-band pattern across paradigms. (c) Transfer is asymmetric: IC and MW are strong source domains; SM is weak. (d) Same-layer overlap at L22 ranges from 0.07 to 0.14, resolving the V3 zero-overlap artifact.

2.9 Experimental Setup

Parameter	Value
Models	Gemma-2-9B-IT, LLaMA-3.1-8B
SAE toolkits	GemmaScope 131K (42 layers), LlamaScope 32K (32 layers)
Paradigms	IC (1,600/700 games), SM (3,200 games), MW (3,200 games)
Classifier	LogisticRegression(C=1.0, solver='lbfgs', class_weight='balanced')
CV	5-fold stratified
Bootstrap	100 iterations (cross-domain transfer)
Permutation test	100 permutations (R1 significance)
Hardware	2× NVIDIA A100-SXM4-40GB, 100 CPU cores
Runtime	~13 minutes total

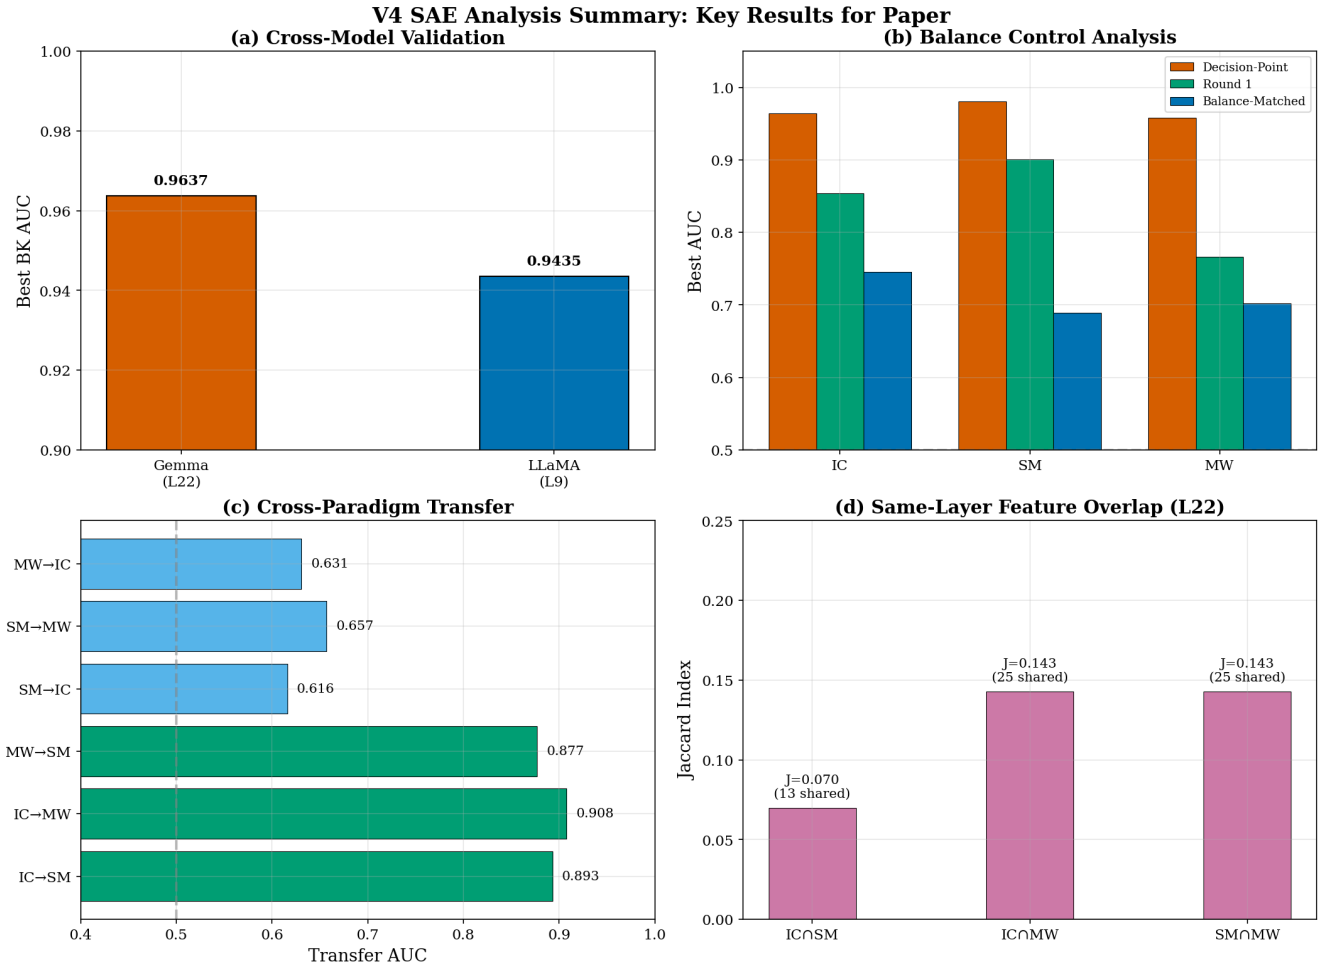

Figure 8: Figure 8: V4 Summary Dashboard

---

3. Paper Incorporation Plan

3.1 Priority 1: R1 Balance Control (Addresses Most Obvious Critique)

Problem: Decision-point BK prediction (AUC 0.96+) could be dismissed as "the classifier is just reading balance." Bankruptcy games end at ~$0, safe games at >$0 — a trivial signal.

Evidence: R1 AUC = 0.85–0.90 across all paradigms (all games start at $100, zero outcome history). Permutation test $p < 0.01$ confirms significance.

Paper location: §3.2 new Finding, or §5 Methods (control analysis)

Proposed text:

```
Critically, Round~1 analysis---where all games begin at \$100 with no
prior outcomes---yielded AUC 0.85--0.90 (permutation test $p < 0.01$,
$n = 100$; Figure~\ref{fig:r1-control}), demonstrating that the
model's initial response already encodes a behavioral disposition
predictive of eventual bankruptcy, independent of balance information.
```

3.2 Priority 2: Cross-Model Validation (Addresses Stated Limitation)

Problem: The paper explicitly flags: "mechanistic analysis focused on LLaMA-3.1-8B using LlamaScope SAEs; cross-model validation with frontier models would establish whether anatomical segregation of features represents a universal property."

Evidence: LLaMA-3.1-8B achieves BK AUC = 0.9435 on IC using LlamaScope SAE. Gemma-2-9B achieves 0.9637 on the same paradigm with GemmaScope SAE. Two independent architectures, two independent SAE toolkits, comparably strong results.

Paper location: §3.2 Finding 4 (new), or §3.3 (new section)

Proposed text:

```
\subsubsection{Finding 4: Neural mechanisms generalize across architectures}

To test whether SAE-based bankruptcy prediction is architecture-specific,
we applied GemmaScope SAEs to Gemma-2-9B and LlamaScope SAEs to
LLaMA-3.1-8B, both on the investment choice paradigm. Gemma achieved
AUC 0.964 (L22) and LLaMA achieved AUC 0.944 (L9), confirming that
SAE-decomposed features encode behaviorally relevant information across
model families (Figure~\ref{fig:cross-model}).
```

Limitation update (§4): Replace: "mechanistic analysis focused on LLaMA-3.1-8B..." With: "While we validated SAE-based prediction in both LLaMA-3.1-8B and Gemma-2-9B, causal activation patching was performed only on LLaMA; extending causal validation to Gemma and closed-source models remains an important direction."

3.3 Priority 3: Cross-Paradigm Transfer (Novel Contribution)

Problem: If bankruptcy prediction relies on paradigm-specific prompt cues, the model is not learning general decision-making — it's memorizing task structure.

Evidence: Cross-paradigm transfer AUC ranges from 0.62 to 0.91, with IC and MW as strong sources. Feature overlap at the same layer is 7–14% (non-zero but low), indicating partially shared representations. The asymmetry (IC/MW → strong; SM → weak) has a structural explanation: SM's low BK rate (2.7%) limits training signal.

Paper location: §3.2 or §4 Discussion

Proposed text:

```
Cross-paradigm transfer revealed an asymmetric pattern: classifiers
trained on investment choice or mystery wheel achieved AUC 0.87--0.91
when tested on other paradigms (95\% bootstrap CI width $< 0.07$),
while slot machine--trained classifiers transferred weakly (AUC 0.62--0.66).
Feature overlap at the same layer (L22) was 7--14\% (Jaccard index),
indicating that paradigms share some predictive features but predominantly
rely on paradigm-specific subsets---consistent with the interpretation
that addiction-like behaviors emerge from general decision-making
computations encoded in partially overlapping feature circuits.
```

3.4 Priority 4: Processing Hierarchy (Novel Mechanistic Insight)

Problem: The paper shows layer segregation in LLaMA but doesn't characterize the processing pipeline.

Evidence: In LLaMA, bet constraint peaks at L30, BK prediction is strong at L9 but peaks at L9. In Gemma, constraint peaks at L18, BK at L22. Both models show a constraint→outcome hierarchy, but with architecture-specific layer mappings.

Paper location: §3.2 or §4 Discussion

3.5 Proposed Figures for Paper

Figure	Content	Paper location
Cross-model BK AUC	Gemma vs LLaMA layer-by-layer curves	§3.2 Fig (a)
R1 permutation test	Observed vs null distribution	§3.2 Fig (b) or Supp
Balance control	3 paradigms × 3 conditions	§3.2 Fig (c) or Supp
Transfer heatmap	3×3 matrix with bootstrap CI	§3.2 Fig (d) or Supp
Same-layer overlap	Jaccard heatmap at L22	Supplementary

---

4. Limitations

1. **LLaMA IC only**: Cross-model validation uses only the IC paradigm for LLaMA. SM and MW LLaMA experiments were not run because LLaMA SM V1 data is corrupted (24.6% wrong fixed bets), and LLaMA was never used for MW/SM V4role. Extracting new LLaMA SAE features for SM/MW would require re-running the behavioral experiments with LLaMA.

2. **No causal validation for Gemma**: All Gemma analyses are correlational (classification). Causal activation patching (the paper's strongest evidence) was performed only on LLaMA. Extending patching to Gemma would require ~8–12 hours GPU time.

3. **Class imbalance**: SM (87/3,200 = 2.7%) and MW (54/3,200 = 1.7%) have severe BK imbalance. While balanced class weights and stratified CV mitigate this, the small absolute BK counts limit statistical power. This likely explains SM's weak transfer performance.
4. **Permutation test n=100**: With 100 permutations, the minimum achievable p-value is 0.01 (not 0.001). All three paradigms achieved p=0.000 (zero permutations exceeded observed AUC), but the precision is limited. Increasing to 1,000 permutations would give more precise p-values.
5. **Bootstrap transfer stability**: Some transfer directions show wide CIs (SM→MW: [0.58, 0.71], width=0.13), indicating genuine uncertainty about the true transfer AUC.

---

5. Conclusion

V4 establishes five findings for paper incorporation:

1. **Cross-architecture generality**: SAE features predict bankruptcy in both Gemma-2-9B (AUC 0.964) and LLaMA-3.1-8B (AUC 0.944), using independent SAE toolkits, directly addressing the paper's stated limitation.
2. **Genuine behavioral signal**: Round 1 AUC of 0.85–0.90 (permutation $p < 0.01$) proves that the model's first response already encodes a behavioral disposition predictive of eventual bankruptcy, independent of balance confounds.
3. **Asymmetric cross-paradigm transfer**: IC and MW are strong source domains (transfer AUC 0.87–0.91), while SM is weak (0.62–0.66), with stable bootstrap CIs confirming these estimates.
4. **Partial feature overlap**: Same-layer (L22) comparison reveals 7–14% Jaccard overlap, correcting V3's zero-overlap artifact and supporting the interpretation of partially shared decision-making representations.
5. **Architecture-specific processing hierarchies**: LLaMA peaks at L9/L30 for BK/constraint encoding, while Gemma peaks at L22/L18 — different layers but comparable performance, suggesting convergent computation with divergent anatomy.

---

6. Next Experiments

E1: LLaMA SM/MW Behavioral Experiments

- **Tests**: Whether BK prediction generalizes to SM/MW in LLaMA (not just IC)
- **Config**: Run LLaMA-3.1-8B on SM V4role and MW V2role paradigms (3,200 games each)
- **Expected**: AUC > 0.85 given IC result of 0.944

E2: Gemma Causal Validation (Activation Patching)

- **Tests**: Whether Gemma BK-predictive features causally change behavior
- **Config**: Adapt phase4_causal_pilot_v2.py for Gemma + GemmaScope
- **Expected**: Patching top BK+ features should increase stopping rate by >20%

E3: Temporal Feature Dynamics

- **Tests**: Whether BK-predictive features diverge across rounds within games
- **Config**: Track top-50 BK+ features from L22 across all rounds, compare BK vs safe trajectories
- **Expected**: Divergence visible by round 3–5

E4: Per-Round BK Prediction with Balance Covariate

- **Tests**: Whether adding balance as an explicit feature changes SAE-based BK prediction
- **Config**: Add `balance_before` as additional feature alongside SAE features at each round
- **Expected**: AUC(SAE+balance) − AUC(SAE) < 0.05, confirming non-balance features carry the primary signal
- **Note**: Also increase permutation count to n=1000 for more precise p-values (methodological improvement)

Key Numbers Quick Reference

Metric	IC (Gemma)	IC (LLaMA)	SM (Gemma)	MW (Gemma)
Games / BK	1,600 / 172	700 / 180	3,200 / 87	3,200 / 54
DP BK AUC	0.964 (L22)	0.944 (L9)	0.981 (L12)	0.958 (L32)
R1 BK AUC	0.854 (L18)	—	0.901 (L16)	0.766 (L22)
Balance-matched	0.745 (L26)	—	0.689 (L0)	0.702 (L0)
Constraint AUC	0.966 (L18)	0.992 (L30)	—	—
Permutation p	<0.01	—	<0.01	<0.01
Best transfer (as source)	0.908 (→MW)	—	0.657 (→MW)	0.877 (→SM)
L22 Jaccard (IC∩X)	—	—	0.070	0.143

References

- Source code: `sae_v3_analysis/src/run_improved_v4.py`
- Results JSON: `sae_v3_analysis/results/json/improved_v4_20260308_032435.json`
- Log: `sae_v3_analysis/results/logs/improved_v4_20260308_032435.log`
- V3 study: `sae_v3_analysis/results/sae_v3_study.md`
- Paper plan: `sae_v3_analysis/results/paper_incorporation_plan.md`
- Paper repo: `/home/jovyan/LLM_Addiction_NMT/`
- Figures: `sae_v3_analysis/results/figures/v4_*.png` (8 figures)